

Module_3_EDA(Exploratory Data Analysis): Understanding the Data

In this module we will perform EDA on student performance dataset 

you can download dataset here https://github.com/kundetivamsi2001/Datasets_AI-ML/blob/main/student_performance.csv
[\(https://github.com/kundetivamsi2001/Datasets_AI-ML/blob/main/student_performance.csv\)](https://github.com/kundetivamsi2001/Datasets_AI-ML/blob/main/student_performance.csv)

In [76]:

```
1 #step 1: Load the libraries dataset
2 import pandas as pd
3 import numpy as np
4 import matplotlib.pyplot as plt
5 import seaborn as sns
6 df=pd.read_csv('student_performance.csv')
7 df.head()
```

Out[76]:

	Student_ID	Age	Gender	Study_Hours	Attendance(%)	Test_Score	Grade
0	S0524	21	Female	2.7	74.0	37.0	F
1	S0603	19	Male	1.5	93.1	26.0	F
2	S0527	20	Female	6.8	97.2	88.0	A
3	S0032	18	Female	6.2	93.5	55.0	C
4	S0617	21	Male	7.4	99.2	77.0	B

2. Have a look on the Data and clean it (if required)

In [29]:

```
1 #2.1 Basic info about dataset
2 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1020 entries, 0 to 1019
Data columns (total 7 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Student_ID      1020 non-null   object 
 1   Age              1020 non-null   int64  
 2   Gender           1020 non-null   object 
 3   Study_Hours     989 non-null   float64
 4   Attendance(%)   990 non-null   float64
 5   Test_Score       1020 non-null   float64
 6   Grade            1011 non-null   object 
dtypes: float64(3), int64(1), object(3)
memory usage: 55.9+ KB
```

```
In [30]: 1 #2.2 check For duplicate values  
2 df.duplicated().sum()
```

Out[30]: 20

```
In [31]: 1 # 2.3 There are 20 duplicate values ,remove them  
2 df=df.drop_duplicates()  
3 df.duplicated().sum()
```

Out[31]: 0

```
In [32]: 1 #2.4 Inspect for missing values and handle them  
2 df.isnull().sum()
```

```
Out[32]: Student_ID      0  
Age            0  
Gender         0  
Study_Hours    30  
Attendance(%) 30  
Test_Score     0  
Grade          8  
dtype: int64
```

```
In [36]: 1 #Fill missing values with mean for Study_Hours with 0,Attendance(%) with m  
2 df['Study_Hours']=df['Study_Hours'].fillna(0)  
3 df['Attendance(%)']=df['Attendance(%)'].fillna(df['Attendance(%)'].mean())  
4 #Fill grade with none  
5 df['Grade']=df['Grade'].fillna('None')  
6 df.isnull().sum()
```

```
Out[36]: Student_ID      0  
Age            0  
Gender         0  
Study_Hours    0  
Attendance(%) 0  
Test_Score     0  
Grade          0  
dtype: int64
```

3. Descriptive Analysis (Summarizing the Numbers)

In [37]: 1 df.describe()

Out[37]:

	Age	Study_Hours	Attendance(%)	Test_Score
count	1000.0000	1000.000000	1000.000000	1000.000000
mean	20.9600	5.006700	84.845773	49.80800
std	2.0036	2.671674	9.820483	21.18745
min	18.0000	-2.000000	50.300000	0.000000
25%	19.0000	3.600000	78.000000	35.000000
50%	21.0000	5.000000	84.845773	50.000000
75%	23.0000	6.300000	91.725000	64.000000
max	24.0000	25.000000	115.600000	100.000000

Interpretation

Example interpretation:

Mean Test_Score = 49.8

Median Test_Score = 50

Insight Some low scorers are pulling the average down.

4. Categorical Data Analysis

In [75]: 1 #check whether the data is balanced or not
2 df["Gender"].value_counts()
3

Out[75]: Male 519
Female 481
Name: Gender, dtype: int64

We have to Check:

Is dataset balanced?

Any bias risk?

```
In [39]: 1 #Grade Distribution
          2 df["Grade"].value_counts()
          3
```

```
Out[39]: F      507
          C      257
          B      148
          A      48
          A+     32
          None    8
          Name: Grade, dtype: int64
```


Insight:

Too many C or D grades may indicate teaching or attendance issues.

5. Visual Analytics: Seeing Patterns

Univariate Analysis

```
In [40]: 1 #5.1 Histogram for Test Score Distribution
          2 plt.hist(df["Test_Score"], bins=10)
          3 plt.xlabel("Test Score")
          4 plt.ylabel("Number of Students")
          5 plt.title("Test Score Distribution")
          6 plt.show()
          7
```



What to look for:

Skewed distribution

Too many low or high scores

Possible cut-off effects


Interpretation: The test scores are normally distributed with mean approx. 50

In [41]:

```

1 #5.2 KDE Plot for Study Hours
2 import seaborn as sns
3 sns.kdeplot(df["Study_Hours"], fill=True)
4 plt.title("Study Hours Density")
5 plt.show()

```



Insight:

Most students study in a narrow range

Very high study hours may not be common

6. Correlation Study

In [42]:

```

1 df.corr()
2


```

Out[42]:

	Age	Study_Hours	Attendance(%)	Test_Score
Age	1.000000	-0.005747	-0.005721	0.024649
Study_Hours	-0.005747	1.000000	0.019061	0.623135
Attendance(%)	-0.005721	0.019061	1.000000	0.012120
Test_Score	0.024649	0.623135	0.012120	1.000000

In [43]:

```
1 #Heatmap
2 plt.figure(figsize=(8,6))
3 sns.heatmap(df.corr(), annot=True, cmap="coolwarm")
4 plt.title("Correlation Heatmap")
5 plt.show()
6
```



Key questions

Does Attendance correlate with Test_Score?

Does Study_Hours matter more than Attendance?

Is Age irrelevant?

7. Scatter Plots (Relationship in Action)


Bivariate Analysis

In [44]:

```

1 plt.scatter(df["Study_Hours"], df["Test_Score"])
2 plt.xlabel("Study Hours")
3 plt.ylabel("Test Score")
4 plt.title("Study Hours vs Performance")
5 plt.show()
6

```



Core pattern (the main story)

There is a strong positive relationship between study hours and test score. This means:

In general, more study hours → better performance.

Outliers and unusual behavior

Now look at the right side:

There are students studying 20–25 hours, but:

Some score high

Some score medium

Some score very low

This tells us:

Conclusion : Studying more does not guarantee high scores.

8. Outlier Exploration (Very Important)

```
In [45]: 1 sns.boxplot(y=df[ "Test_Score" ])
2 plt.title("Outliers in Test Score")
3 plt.show()
```



9. Feature Interactions (EDA Level-Up)

```
In [46]: 1 #Grade vs Attendance
2 df.groupby("Grade")[ "Attendance(%)"].mean()
```

```
Out[46]: Grade
A      87.364231
A+     85.450000
B      84.882685
C      84.923360
F      84.551379
None   82.800000
Name: Attendance(%), dtype: float64
```

Insight:

Higher grades usually have higher attendance.

```
In [47]: 1 #Gender vs Performance
2 df.groupby("Gender")[ "Test_Score"].mean()
```

```
Out[47]: Gender
Female    48.717256
Male      50.818882
Name: Test_Score, dtype: float64
```

Conclusion : Male Category have high scores

10. Create Buckets (creating new feature)

In [48]:

```

1 #Attendance Buckets
2 df["Attendance_Level"] = pd.cut(df["Attendance(%)"],bins=[0, 60, 80, 100],
3                                 labels=["Low", "Medium", "High"])
4 )
5 df.head()
6

```

Out[48]:

	Student_ID	Age	Gender	Study_Hours	Attendance(%)	Test_Score	Grade	Attendance_Level
0	S0524	21	Female	2.7	74.0	37.0	F	Medium
1	S0603	19	Male	1.5	93.1	26.0	F	High
2	S0527	20	Female	6.8	97.2	88.0	A	High
3	S0032	18	Female	6.2	93.5	55.0	C	High
4	S0617	21	Male	7.4	99.2	77.0	B	High

In [49]:

```
1 df.groupby("Attendance_Level")["Test_Score"].mean()
```

Out[49]:

```

Attendance_Level
Low      44.200000
Medium   49.859016
High     49.794913
Name: Test_Score, dtype: float64


```

We have created new feature called Attendance_level and analysed

11 Univariate Analysis on Numerical Variables

In [50]:

```
1 #visualize
2 plt.hist(df["Age"], bins=10)
3 plt.xlabel("Age")
4 plt.ylabel("Number of Students")
5 plt.title("Age Distribution")
6 plt.show()
```



In []:

1